

Concurrency in Scala

Mikhail Mutcianko, Alexey Shcherbakov

СПБГУ, СП

8 апреля 2021

Concurrent, Parallel, Asynchronous

Concurrent vs Parallel computations

Concurrent

Concurrent computing is a form of computing in which several computations are executed during overlapping time periods — instead of sequentially, with one completing before the next.

Parallel


Parallel computations must also advance simultaneously

Concurrent vs Parallel computations

Concepts in Concurrency



Concurrent, non-parallel execution



Concurrent, parallel execution

Synchronous execution

In synchronous model of execution for any statement S of a linear top-down program flow, computation of all statements that are defined before S must finish by the time S is started.

Synchronous execution

In synchronous model of execution for any statement S of a linear top-down program flow, computation of all statements that are defined before S must finish by the time S is started.

```
1 | val a = longComputeA()  
2 | val b = longComputeB() // a is initialized  
3 |  
4 | longComputeAB(a,b) // both a and b are computed
```

Asynchronous execution

Asynchronous computations drop the «happens before» requirement. Async model allows to continue the next statement even if the previous has not finished yet.

Asynchronous execution

```
1 | println("starting A")
2 | val a = Future { longComputeA(); println("A finished") }
3 | println("starting B")
4 | val b = Future { longComputeB(); println("B finished") }
5 | println("starting AB")
6 | longComputeAB(a,b)
```


Asynchronous execution

```
1 println("starting A")
2 val a = Future { longComputeA(); println("A finished") }
3 println("starting B")
4 val b = Future { longComputeB(); println("B finished") }
5 println("starting AB")
6 longComputeAB(a,b)
```

```
starting A
starting B
starting AB
A finished
B finished
```

Concurrency Primitives

Runnable and Callable

Low-level computation abstractions are defined as traits

```
1 // computation does NOT return a value
2 trait Runnable {
3     def run(): Unit
4 }
5
6 // computation returns a value
7 trait Callable[V] {
8     def call(): V
9 }
```

Thread

Scala concurrency is built on top of the Java concurrency model.

A Thread takes a Runnable. You have to call start on a Thread in order for it to run the Runnable.

```
1  val hello = new Thread(new Runnable {  
2      def run() {  
3          println("hello world")  
4      }  
5  })  
6  
7  hello.start  
8  // -> hello world
```

Thread

Scala concurrency is built on top of the Java concurrency model.

A Thread takes a Runnable. You have to call start on a Thread in order for it to run the Runnable.

```
1 | val hello = new Thread(() => println("hello world"))
2 |
3 | hello.start
4 | // -> hello world
```

Futures

Future is an abstraction capturing a computation process

A computation can be in one of the following **three** states:

- unfinished
- successful
- failed

Future Methods

Common

```
| def Future.apply(f: =>A): Future[A]
```

- accepts a code block
- starts computation eagerly

Future Methods

Common

```
| def Future.apply(f: =>A): Future[A]
```

- accepts a code block
- starts computation eagerly

```
| def onComplete(cb: Try[A] =>B): Unit
```

- provides a callback for a finished computation

Future Methods

Common

```
| def Future.apply(f: =>A)(implicit ec: ExecutionContext): Future[A]
```

- accepts a code block
- starts computation eagerly

```
| def onComplete(cb: Try[A] =>B)(implicit ec: ExecutionContext): Unit
```

- provides a callback for a finished computation

ExecutionContext

An *ExecutionContext* is an abstraction which implementations must provide a concrete mechanism of executing computations represented by Scala's Future. The intent of ExecutionContext is to lexically scope code execution

- built-in implementation:

```
scala.concurrent.ExecutionContext.Implicits.global
```

- based on Java's ForkJoinPool
- automatically scales to the number of CPU cores

Working with results

There are several ways of processing an async computation's result:

- block current thread and get a result
- execute a callback when a computation finishes
- compose another computation

Blocking the result

An asynchronous computation can be turned into a synchronous one by using `Await.result`

This method will block the current thread until:

- the Future completes(successfully or not)
- a timeout occurs

Blocking the result

Example

```
1 val greetingFuture = Future {  
2     Thread.sleep(1000)  
3     "Hello"  
4 }  
5  
6 val greeting: String = Await.result(greetingFuture, Duration.Inf)  
7  
8 println(s"Result: $greeting")
```

Blocking the result

Poor practice

Blocking on a future is strongly discouraged for the sake of performance and for the prevention of deadlocks. Callbacks and combinators on futures are a preferred way to use their results.

Blocking the result

Deadlock example

```
1  implicit val ec = ExecutionContext
2    .fromExecutor(Executors.newFixedThreadPool(1))
3
4  def addOne(x: Int) = Future(x + 1)
5
6  def multiply(x: Int, y: Int) = Future {
7    val a = addOne(x)
8    val b = addOne(y)
9    val result = for (r1 <- a; r2 <- b) yield r1 * r2
10
11    // this will dead-lock
12    Await.result(result, Duration.Inf)
13  }
```


Using callbacks

onComplete

Processing the result of a Future can be done in a completely non-blocking way by providing a callback using: `Future.onComplete[U](f: (Try[T]) => U): Unit`

```
1 val f: Future[List[String]] = Future {
2   session.getRecentPosts
3 }
4
5 f onComplete {
6   case Success(posts) => for (post <- posts) println(post)
7   case Failure(t) => println("An error has occurred: " + t.getMessage)
8 }
```

Using callbacks

foreach

In the case where only successful results need to be handled, the `foreach` callback can be used:

```
1  val f: Future[List[String]] = Future {  
2      session.getRecentPosts  
3  }  
4  
5  f foreach { posts =>  
6      for (post <- posts) println(post)  
7  }
```

Callback hell

Chaining computations using callbacks is achieved by nesting creation of Futures. This can lead to severe cases of LOP - Ladder Oriented Programming also known as «callback hell»

Callback hell

Example

```
1  queryDb(8612).onComplete {
2      case Failure(ex: Exception) =>
3          println(s"Operation failed with $ex")
4      case Success(fileName: String) =>
5          loadFileAsync(fileName).onComplete {
6              case Failure(ex: Exception) =>
7                  println(s"Operation failed with $ex")
8              case Success(url: String) =>
9                  loadPageAsync(url).onComplete {
10                      case Failure(ex: Exception) => println(s"Operation failed with $ex")
11                      case Success(text: String) => Future { ... }
12                  ...
13      }
```

Functional composition of Futures

Monad

In Scala a `Future[+A]` is a monad, providing the following methods:

- `def map[B](f: A => B): Future[B]`
- `def flatMap[B](f: A => Future[B]): Future[B]`
- `def withFilter(f: A => Boolean): Future[A]`

Functional composition of Futures

Example

Monadic composition style:

```
1 | queryDb(8612)
2 |   .flatMap(fileName => loadFileAsync(fileName))
3 |   .flatMap(url      => loadPageAsync(url))
4 |   .flatMap(pageText => println(pageText))
```

Functional composition of Futures

Example

Monadic composition style:

```
1 | queryDb(8612)
2 |   .flatMap(fileName => loadFileAsync(fileName))
3 |   .flatMap(url      => loadPageAsync(url))
4 |   .flatMap(pageText => println(pageText))
```

For-comprehension style:

```
1 | for {
2 |   fileName <- queryDb(8612)
3 |   url      <- loadFileAsync(fileName)
4 |   pageText <- loadPageAsync(url)
5 | } println(pageText)
```

Error handling

- `def recover[U >: T](pf: PartialFunction[Throwable, U]): Future[U]`
Creates a new future that will handle any matching throwable that this future might contain.
- `def recoverWith[U >: T](pf: PartialFunction[Throwable, Future[U]]): Future[U]`
Same as `recover`, but composes another future instead of a value
- `def fallbackTo[U >: T](that: Future[U]): Future[U]`
Creates a new future which holds the result of this future if it was completed successfully, or, if not, the result of the `that` future if that is completed successfully.
- `def failed: Future[Throwable]`
Returns a Future with an exception as a result value if the original one has failed

Error handling

Example

```
1 Future (6 / 0) recover { case e: ArithmeticException => 0 } // result: 0
2
3 val f = Future { Int.MaxValue }
4 Future (6 / 0) recoverWith { case e: ArithmeticException => f } // result: Int.MaxValue
5
6 val f = Future { throw new RuntimeException("failed") }
7 val g = Future { 5 }
8 val h = f fallbackTo g
9 h foreach println // Eventually prints 5
10
11 val g = Future { 2 / 0 }
12 for (exc <- g.failed) println(exc) // result: java.lang.ArithmeticException: / by zero
```

Future aggregation

- `def traverse(in: M[A])(fn: (A) => Future[B]): Future[M[B]]`

Asynchronously and non-blockingly transforms a `IterableOnce[A]` into a `Future[IterableOnce[B]]` using the provided function `A => Future[B]`

- `def sequence(in: M[Future[B]]): Future[M[B]]`

Transforms a sequence of futures into a future of sequences

- `def zip(other: Future[B]): Future[(A, B)]`

Creates a single future of `Tuple2` from two futures

- `def foldLeft(futures: Iterable[Future[T]])(zero: R)(op: (R, T) => R): Future[R]`

- `def reduceLeft(futures: Iterable[Future[T]])(op: (R, T) => R): Future[R]`

Promise

Promise is an API for creating Futures with a controllable state. Promises *complete* the Futures they produce (by "completing" the promise)

The generated Future state can be controlled with the following:

- `complete / completeWith`
- `tryComplete / tryCompleteWith`
- `success / trySuccess / failure / tryFailure`

Promise

Example

```
1  val p = Promise[T]()
2  val f = p.future
3
4  val producer = Future {
5    val r = produceSomething()
6    p success r
7    continueDoingSomethingUnrelated()
8  }
9
10 val consumer = Future {
11   startDoingSomething()
12   f foreach { r => doSomethingWithResult() }
13 }
```

Promise assignment semantics

Promises have single-assignment semantics. As such, they can be completed only once.

Calling success on a promise that has already been completed (or failed) will throw an `IllegalStateException`

Parallel Collections

Scala parallel collections

Scala provides an easy way of converting any sequential collection to parallel with `.par`

Scala parallel collections

Scala provides an easy way of converting any sequential collection to parallel with `.par`

```
1 | val list = (1 to 10000).toList  
2 | list.par.map(_ + 42)
```

“out-of-order” semantics of parallel collections lead to the following two implications:

- Side-effecting operations can lead to non-determinism
- Non-associative operations lead to non-determinism

Practice